

499-hw1-part3

Batuhan Karaca 2310191

November 2022

List of plots

1	Accuracy values for 0.0001, 30, 3, 128, ReLU(), 463	6
2	Loss values for 0.0001, 30, 3, 128, ReLU(), 463	7
3	Accuracy values for 0.0001, 30, 3, 128, ReLU(), 50004	8
4	Loss values for 0.0001, 30, 3, 128, ReLU(), 50004	9
5	Accuracy values for 0.0001, 30, 3, 128, Sigmoid(), 463	10
6	Loss values for 0.0001, 30, 3, 128, Sigmoid(), 463	11
7	Accuracy values for 0.0001, 30, 3, 128, Sigmoid(), 50004	12
8	Loss values for 0.0001, 30, 3, 128, Sigmoid(), 50004	13
9	Accuracy values for 0.0001, 30, 3, 128, Tanh(), 463	14
10	Loss values for 0.0001, 30, 3, 128, Tanh(), 463	15
11	Accuracy values for 0.0001, 30, 3, 128, Tanh(), 50004	16
12	Loss values for 0.0001, 30, 3, 128, Tanh(), 50004	17
13	Accuracy values for 0.0001, 30, 3, 16, ReLU(), 463	18
14	Loss values for 0.0001, 30, 3, 16, ReLU(), 463	19
15	Accuracy values for 0.0001, 30, 3, 16, ReLU(), 50004	20
16	Loss values for 0.0001, 30, 3, 16, ReLU(), 50004	21
17	Accuracy values for 0.0001, 30, 3, 16, Sigmoid(), 463	22
18	Loss values for 0.0001, 30, 3, 16, Sigmoid(), 463	23
19	Accuracy values for 0.0001, 30, 3, 16, Sigmoid(), 50004	24
20	Loss values for 0.0001, 30, 3, 16, Sigmoid(), 50004	25
21	Accuracy values for 0.0001, 30, 3, 16, Tanh(), 463	26
22	Loss values for 0.0001, 30, 3, 16, Tanh(), 463	27
23	Accuracy values for 0.0001, 30, 3, 16, Tanh(), 50004	28
24	Loss values for 0.0001, 30, 3, 16, Tanh(), 50004	29
25	Accuracy values for 0.01, 30, 3, 128, ReLU(), 463	30
26	Loss values for 0.01, 30, 3, 128, ReLU(), 463	31
27	Accuracy values for 0.01, 30, 3, 128, ReLU(), 50004	32
28	Loss values for 0.01, 30, 3, 128, ReLU(), 50004	33
29	Accuracy values for 0.01, 30, 3, 128, Sigmoid(), 463	34
30	Loss values for 0.01, 30, 3, 128, Sigmoid(), 463	35
31	Accuracy values for 0.01, 30, 3, 128, Sigmoid(), 50004	36
32	Loss values for 0.01, 30, 3, 128, Sigmoid(), 50004	37

33	Accuracy values for 0.01, 30, 3, 128, Tanh(), 463	38
34	Loss values for 0.01, 30, 3, 128, Tanh(), 463	39
35	Accuracy values for 0.01, 30, 3, 128, Tanh(), 50004	40
36	Loss values for 0.01, 30, 3, 128, Tanh(), 50004	41
37	Accuracy values for 0.01, 30, 3, 16, ReLU(), 463	42
38	Loss values for 0.01, 30, 3, 16, ReLU(), 463	43
39	Accuracy values for 0.01, 30, 3, 16, ReLU(), 50004	44
40	Loss values for 0.01, 30, 3, 16, ReLU(), 50004	45
41	Accuracy values for 0.01, 30, 3, 16, Sigmoid(), 463	46
42	Loss values for 0.01, 30, 3, 16, Sigmoid(), 463	47
43	Accuracy values for 0.01, 30, 3, 16, Sigmoid(), 50004	48
44	Loss values for 0.01, 30, 3, 16, Sigmoid(), 50004	49
45	Accuracy values for 0.01, 30, 3, 16, Tanh(), 463	50
46	Loss values for 0.01, 30, 3, 16, Tanh(), 463	51
47	Accuracy values for 0.01, 30, 3, 16, Tanh(), 50004	52
48	Loss values for 0.01, 30, 3, 16, Tanh(), 50004	53
49	Accuracy values for 1, 30, 3, 128, ReLU(), 463	54
50	Loss values for 1, 30, 3, 128, ReLU(), 463	55
51	Accuracy values for 1, 30, 3, 128, ReLU(), 50004	56
52	Loss values for 1, 30, 3, 128, ReLU(), 50004	57
53	Accuracy values for 1, 30, 3, 128, Sigmoid(), 463	58
54	Loss values for 1, 30, 3, 128, Sigmoid(), 463	59
55	Accuracy values for 1, 30, 3, 128, Sigmoid(), 50004	60
56	Loss values for 1, 30, 3, 128, Sigmoid(), 50004	61
57	Accuracy values for 1, 30, 3, 128, Tanh(), 463	62
58	Loss values for 1, 30, 3, 128, Tanh(), 463	63
59	Accuracy values for 1, 30, 3, 128, Tanh(), 50004	64
60	Loss values for 1, 30, 3, 128, Tanh(), 50004	65
61	Accuracy values for 1, 30, 3, 16, ReLU(), 463	66
62	Loss values for 1, 30, 3, 16, ReLU(), 463	67
63	Accuracy values for 1, 30, 3, 16, ReLU(), 50004	68
64	Loss values for 1, 30, 3, 16, ReLU(), 50004	69
65	Accuracy values for 1, 30, 3, 16, Sigmoid(), 463	70
66	Loss values for 1, 30, 3, 16, Sigmoid(), 463	71
67	Accuracy values for 1, 30, 3, 16, Sigmoid(), 50004	72
68	Loss values for 1, 30, 3, 16, Sigmoid(), 50004	73
69	Accuracy values for 1, 30, 3, 16, Tanh(), 463	74
70	Loss values for 1, 30, 3, 16, Tanh(), 463	75
71	Accuracy values for 1, 30, 3, 16, Tanh(), 50004	76
72	Loss values for 1, 30, 3, 16, Tanh(), 50004	77

1 Introduction

Note that the terms *combination* and *configuration* are used interchangeably.

This report aims to analyze the results obtained by a model on MNIST dataset using different combination of hyperparameters. The methodology is as follows. For the preprocessing stage, the pixel values are normalized to the values in the range $[0,1]$. Weighted pixel values would give another value in the same range. Tanh or sigmoid would output a value in a smaller subset of range $[0, \sim 0.7]$. In this manner, the output values won't explode avoiding the precision loss. The data is sampled using a shuffled sample (**shuffle=True** in **DataLoader**) to achieve more generalized results, avoiding overfitting. In the first step, with a grid parameter search the model is trained with each configuration on a train set, tested on a validation set. In the second step the configuration with the highest accuracy score is tested with another independent test set. This time the train and validation sets are combined into a new train set. The test set needs to be independent in order to avoid overfitting.

In the next section (section 2), the tested hyperparameters are introduced with the decisions made for the values preferred for each. Then, the results for the first step are analyzed for the parameter. There comes the result of the second step after that (section 3). In the Appendix section, there is a loss and accuracy graph for each combination. Each graph shows the epochs in its x-axis whereas the corresponding loss or accuracy value in its y-axis. The combinations can be seen in the figure captions. There are 36 combinations in total, as 3, 1, 1, 2, 3 and 2 different values are tested for learning rate, number of epochs, depth, breadth, activation function and batch size respectively.

2 First step

GPU is used when training the network. The hyperparameters are chosen accordingly. There are 6 hyperparameters tested in this experiment.

First and foremost, 3 different values of learning rates are evaluated, 0.0001, 0.01 and 1. The values are purely a matter of choice that no reasoning was made. As the results show, 0.0001 was so small that not enough advancements have been made in the given epochs. On the other hand, the value 1 was so large that the updates were unstable. The model is moving so uncontrollably and rapidly in the loss space. This divergent behavior resulted in model not converging, possibly needing more training epochs. Thus, the model performed poorly for this value. 0.01 gave the best results among the values.

Secondly, another main parameter, number of epochs, comes into consideration. The only value tested was 30. Reducing the execution time is a reason for why a single value is preferred. Furthermore, as seen in the plots (Appendix section), 30 is redundant for some of the combinations. The model skyrockets to about 99% in as small as 10 epochs for the combination 0.01, 30, 3, 128, any activation function, 463. We can say that 30 was a sufficient value for this task. Also note that as opposed to the theory that every batch of information is sampled in one epoch, we are iterating over every batches in a single epoch. Therefore, the number of epochs is relatively small compared to hundreds and thousands. Note that overfitting starts to be observed when the test loss is

minimum, or test accuracy is maximum. After that point, we would see the test loss or accuracy diverging from the train loss or accuracy, and the model starts to perform better when training, more poorly when testing. However, SGD or Mini-batch gradient descent's wavy behaviour makes finding this point more difficult. We can avoid overfitting using a regularization technique such as dropout, or batch normalization. Regularization was not used in the experiment.

Note that larger or smaller number of parameters do not necessarily increase the accuracy. Insufficient number of parameters would underfit needing more training time, whereas redundant parameters may overfit the model. Nevertheless, on an interval the accuracy keeps climbing as there are more parameters that can fit. To avoid overfitting, we need to set our breadth and depth values accordingly.

Thirdly, depth parameter (Number of hidden parameters + 1) is considered. Larger numbers are not considered to reduce the execution time and avoid overfitting. Compared to the number of parameters in a layer, adding layers requires sequential computation, hence more matrix multiplications, slowing down the execution. It is observed in the results that the number 3 is sufficient for this task. The choice is not purely coincidental as I came across an implementation successful with the MNIST dataset using the number 3. The video showing the implementation is in the YouTube link below; however, it is not the original source. The original author is Mike Cohen, giving the Udemy course *A deep understanding of deep learning (with Python intro)* [1].

<https://www.youtube.com/watch?v=RLdxbvMZTsg>

Numbers 16 and 128 were chosen for the breadth (number of hidden parameters) parameter. I wanted to see if there is a difference between the values as small as 16 and as large as 128, and if 16 would perform well. I did not prefer to use larger values. Large parameters may mean more resources on the RAM (and VRAM for the GPU). Generally, increasing breadth value increases the accuracy. Leaving other parameters equal, only increasing the breadth parameter gives small to large improvements in the accuracy results. However the model performs depending on the other parameters as well. For the configuration 0.01, 30, 3, 16, *Sigmoid()*; the model performed slightly worse than its counterpart with 128 nodes (about 97 to 99). Hence, depending on the other hyperparameters, the value 16 either underfits or does not underfit.

In addition to the Sigmoid and Tanh functions, the commonly used ReLU(Rectified linear unit) is used. It solves the vanishing gradient and precision loss problems of both functions, and computationally less expensive. However, the results does not signify this. The reason could be that the vanishing problem is more visible with deeper networks. Because in the backpropagation step, as gradient is getting smaller, deeper networks have their earlier layers' gradients more likely to be zero. The value 3 seems to not let the observation of this phenomenon. Further note that if the pixel values are not normalized, ReLU should perform better than the other two, avoiding the preprocessing stage, reducing the computation cost more.

One last parameter, the batch size, is considered. Theoretically, the Stochas-

tic gradient descent (SGD) takes one sample at one epoch. This makes it more advantageous to the Batch gradient descent as it updates its parameters more frequently for the same amount of data. It can converge faster for large data. The instability of SGD may even make it find a better minima. However, the GPU cannot parallelize a single sample at a time and as our implementation requires, all the data should be scanned in a single epoch, slowing down the execution. The batch size is compromised with 463 samples per batch, removing 1 samples per batch (SGD) from the list for time concerns. The results show that for the batch size of 50004 (All the train data) the model converges monotonously, whereas for batch size of 463 the model is fluctuating. However, compared to the learning rate, the model is not diverging, instead converging and fluctuating around the minima. The results indicate that the batch size of 463 performs better than that of 50004. Mini-batch gradient descent still converges more earlier than the Batch gradient descent similar to SGD, also benefiting from the GPU.

Learning rate	Number of epochs	Depth	Breadth	Activation function	Batch size	Test accuracy
0.0001	30	3	16	<i>ReLU()</i>	463	90.28077697753906 \pm 0.6855044960975647
0.0001	30	3	16	<i>ReLU()</i>	50004	13.48552131652832 \pm 1.7058191299438477
0.0001	30	3	16	<i>Sigmoid()</i>	463	65.70194244384766 \pm 4.243033409118652
0.0001	30	3	16	<i>Sigmoid()</i>	50004	10.742541313171387 \pm 0.889334499835968
0.0001	30	3	16	<i>Tanh()</i>	463	90.53995513916016 \pm 1.242717981338501
0.0001	30	3	16	<i>Tanh()</i>	50004	19.584632873535156 \pm 3.5564305782318115
0.0001	30	3	128	<i>ReLU()</i>	463	95.50755310058594 \pm 0.684922456741333
0.0001	30	3	128	<i>ReLU()</i>	50004	45.7717399597168 \pm 5.663984775543213
0.0001	30	3	128	<i>Sigmoid()</i>	463	91.5982666015625 \pm 0.5833438634872437
0.0001	30	3	128	<i>Sigmoid()</i>	50004	14.083272933959961 \pm 2.4819447994232178
0.0001	30	3	128	<i>Tanh()</i>	463	95.20518493652344 \pm 0.35586246848106384
0.0001	30	3	128	<i>Tanh()</i>	50004	62.70338821411133 \pm 2.0753366947174072
0.01	30	3	16	<i>ReLU()</i>	463	96.95463562011719 \pm 0.5730116367340088
0.01	30	3	16	<i>ReLU()</i>	50004	87.30882263183594 \pm 0.6131219267845154
0.01	30	3	16	<i>Sigmoid()</i>	463	97.86177062988281 \pm 0.430091917514801
0.01	30	3	16	<i>Sigmoid()</i>	50004	58.00975799560547 \pm 2.31934785428955
0.01	30	3	16	<i>Tanh()</i>	463	97.55940246582031 \pm 0.520574152469635
0.01	30	3	16	<i>Tanh()</i>	50004	87.34381866455078 \pm 0.49962684512138367
0.01	30	3	128	<i>ReLU()</i>	463	99.43843078613281 \pm 0.2292753905057907
0.01	30	3	128	<i>ReLU()</i>	50004	93.43932342529297 \pm 0.1511632800102234
0.01	30	3	128	<i>Sigmoid()</i>	463	99.95680236816406 \pm 0.08466586470603943
0.01	30	3	128	<i>Sigmoid()</i>	50004	88.25013732910156 \pm 0.18505258858203888
0.01	30	3	128	<i>Tanh()</i>	463	99.30885314941406 \pm 0.24277380108833313
0.01	30	3	128	<i>Tanh()</i>	50004	94.85101318359375 \pm 0.11484114825725555
1	30	3	16	<i>ReLU()</i>	463	9.740819931030273 \pm 0.8777143955230713
1	30	3	16	<i>ReLU()</i>	50004	10.315975189208984 \pm 0.3357134163379669
1	30	3	16	<i>Sigmoid()</i>	463	9.244059562683105 \pm 1.099924087524414
1	30	3	16	<i>Sigmoid()</i>	50004	11.169106483459473 \pm 1.9723540544509888
1	30	3	16	<i>Tanh()</i>	463	14.600430488586426 \pm 3.8624627590179443
1	30	3	16	<i>Tanh()</i>	50004	42.551597595214844 \pm 7.36905574798584
1	30	3	128	<i>ReLU()</i>	463	10.820734024047852 \pm 0.9327060580253601
1	30	3	128	<i>ReLU()</i>	50004	10.439764976501465 \pm 0.7678419947624207
1	30	3	128	<i>Sigmoid()</i>	463	9.438444137573242 \pm 0.7414588928222656
1	30	3	128	<i>Sigmoid()</i>	50004	9.80341625213623 \pm 0.4264678359031677
1	30	3	128	<i>Tanh()</i>	463	9.913606643676758 \pm 0.9241273999214172
1	30	3	128	<i>Tanh()</i>	50004	69.65682983398438 \pm 2.995284080505371

Table 1: The results for each combination

3 Second step

The configuration with the highest accuracy score is 0.01, 30, 3, 128, *Sigmoid()*, 463. The new accuracy score is 100.0 ± 0.0 .

4 Appendix




Figure 1: Accuracy values for 0.0001, 30, 3, 128, ReLU(), 463




Figure 2: Loss values for 0.0001, 30, 3, 128, ReLU(), 463




Figure 3: Accuracy values for 0.0001, 30, 3, 128, ReLU(), 50004




Figure 4: Loss values for 0.0001, 30, 3, 128, ReLU(), 50004




Figure 5: Accuracy values for 0.0001, 30, 3, 128, Sigmoid(), 463




Figure 6: Loss values for 0.0001, 30, 3, 128, Sigmoid(), 463




Figure 7: Accuracy values for 0.0001, 30, 3, 128, Sigmoid(), 50004




Figure 8: Loss values for 0.0001, 30, 3, 128, Sigmoid(), 50004




Figure 9: Accuracy values for 0.0001, 30, 3, 128, Tanh(), 463




Figure 10: Loss values for 0.0001, 30, 3, 128, Tanh(), 463




Figure 11: Accuracy values for 0.0001, 30, 3, 128, Tanh(), 50004




Figure 12: Loss values for 0.0001, 30, 3, 128, Tanh(), 50004




Figure 13: Accuracy values for 0.0001, 30, 3, 16, ReLU(), 463




Figure 14: Loss values for 0.0001, 30, 3, 16, ReLU(), 463




Figure 15: Accuracy values for 0.0001, 30, 3, 16, ReLU(), 50004




Figure 16: Loss values for 0.0001, 30, 3, 16, ReLU(), 50004




Figure 17: Accuracy values for 0.0001, 30, 3, 16, Sigmoid(), 463




Figure 18: Loss values for 0.0001, 30, 3, 16, Sigmoid(), 463




Figure 19: Accuracy values for 0.0001, 30, 3, 16, Sigmoid(), 50004




Figure 20: Loss values for 0.0001, 30, 3, 16, Sigmoid(), 50004




Figure 21: Accuracy values for 0.0001, 30, 3, 16, Tanh(), 463




Figure 22: Loss values for 0.0001, 30, 3, 16, Tanh(), 463




Figure 23: Accuracy values for 0.0001, 30, 3, 16, Tanh(), 50004




Figure 24: Loss values for 0.0001, 30, 3, 16, Tanh(), 50004




Figure 25: Accuracy values for 0.01, 30, 3, 128, ReLU(), 463




Figure 26: Loss values for 0.01, 30, 3, 128, ReLU(), 463




Figure 27: Accuracy values for 0.01, 30, 3, 128, ReLU(), 50004




Figure 28: Loss values for 0.01, 30, 3, 128, ReLU(), 50004




Figure 29: Accuracy values for 0.01, 30, 3, 128, Sigmoid(), 463




Figure 30: Loss values for 0.01, 30, 3, 128, Sigmoid(), 463




Figure 31: Accuracy values for 0.01, 30, 3, 128, Sigmoid(), 50004




Figure 32: Loss values for 0.01, 30, 3, 128, Sigmoid(), 50004




Figure 33: Accuracy values for 0.01, 30, 3, 128, Tanh(), 463




Figure 34: Loss values for 0.01, 30, 3, 128, Tanh(), 463




Figure 35: Accuracy values for 0.01, 30, 3, 128, Tanh(), 50004




Figure 36: Loss values for 0.01, 30, 3, 128, Tanh(), 50004




Figure 37: Accuracy values for 0.01, 30, 3, 16, ReLU(), 463




Figure 38: Loss values for 0.01, 30, 3, 16, ReLU(), 463




Figure 39: Accuracy values for 0.01, 30, 3, 16, ReLU(), 50004




Figure 40: Loss values for 0.01, 30, 3, 16, ReLU(), 50004




Figure 41: Accuracy values for 0.01, 30, 3, 16, Sigmoid(), 463




Figure 42: Loss values for 0.01, 30, 3, 16, Sigmoid(), 463




Figure 43: Accuracy values for 0.01, 30, 3, 16, Sigmoid(), 50004




Figure 44: Loss values for 0.01, 30, 3, 16, Sigmoid(), 50004




Figure 45: Accuracy values for 0.01, 30, 3, 16, Tanh(), 463




Figure 46: Loss values for 0.01, 30, 3, 16, Tanh(), 463




Figure 47: Accuracy values for 0.01, 30, 3, 16, Tanh(), 50004




Figure 48: Loss values for 0.01, 30, 3, 16, Tanh(), 50004




Figure 49: Accuracy values for 1, 30, 3, 128, ReLU(), 463




Figure 50: Loss values for 1, 30, 3, 128, ReLU(), 463




Figure 51: Accuracy values for 1, 30, 3, 128, ReLU(), 50004




Figure 52: Loss values for 1, 30, 3, 128, ReLU(), 50004




Figure 53: Accuracy values for 1, 30, 3, 128, Sigmoid(), 463




Figure 54: Loss values for 1, 30, 3, 128, Sigmoid(), 463




Figure 55: Accuracy values for 1, 30, 3, 128, Sigmoid(), 50004




Figure 56: Loss values for 1, 30, 3, 128, Sigmoid(), 50004




Figure 57: Accuracy values for 1, 30, 3, 128, Tanh(), 463




Figure 58: Loss values for 1, 30, 3, 128, Tanh(), 463




Figure 59: Accuracy values for 1, 30, 3, 128, Tanh(), 50004




Figure 60: Loss values for 1, 30, 3, 128, Tanh(), 50004




Figure 61: Accuracy values for 1, 30, 3, 16, ReLU(), 463




Figure 62: Loss values for 1, 30, 3, 16, ReLU(), 463




Figure 63: Accuracy values for 1, 30, 3, 16, ReLU(), 50004




Figure 64: Loss values for 1, 30, 3, 16, ReLU(), 50004




Figure 65: Accuracy values for 1, 30, 3, 16, Sigmoid(), 463




Figure 66: Loss values for 1, 30, 3, 16, Sigmoid(), 463




Figure 67: Accuracy values for 1, 30, 3, 16, Sigmoid(), 50004




Figure 68: Loss values for 1, 30, 3, 16, Sigmoid(), 50004




Figure 69: Accuracy values for 1, 30, 3, 16, Tanh(), 463




Figure 70: Loss values for 1, 30, 3, 16, Tanh(), 463




Figure 71: Accuracy values for 1, 30, 3, 16, Tanh(), 50004




Figure 72: Loss values for 1, 30, 3, 16, Tanh(), 50004

5 References

References

- [1] Cohen, M. X. (2022, October). A deep understanding of deep learning (with python intro). Udemy. Retrieved November 6, 2022, from https://www.udemy.com/course/deeplearning_x